

# diags: R Tools for Catch per Unit Effort Analysis

Laurence Kell

ICCAT

---

## Abstract

The **diags** package provides methods for plotting and summarising Catch per Unit Effort (CPUE) data used in fitting fish stock assessment methods. Programs for stock assessment are generally implemented as standalone executable programs with their own text files for input and output files. **diags** provides a set of methods for reading these data and the results from fitting them into R.

*Keywords:* R, CPUE, diagnostics, residuals, stock assessment.

---

## Contents

## 1. introduction

The **diags** package provides methods for plotting and summarising Catch per Unit Effort (CPUE) data used in fitting fish stock assessment methods. Programs for stock assessment are generally implemented as standalone executable programs with their own text files for input and output files. **diags** provides a set of methods for reading these data and the results from fitting them into R.

The stock assessment package for which there are R methods for reading text files are

- **ASPIC** a biomass dynamic model fitted by maximising the likelihood
- **BSP** a Bayesian biomass dynamic model fitted using the SIR algorithm
- **VPA Suite** Input file format mainly used by ICES for virtual population analysis
- **VPA2Box** An age structured model based on virtual population analysis
- **Multifan-CL** A statistical, length-based, age-structured model
- **Stock Synthesis** age and size structure assessment model

## 2. Data

The `readCpue` method reads data from the various stock assessment files into a common data frame. There is an example data frame in the package

```
> library(diags)
> data(rsd1)
> head(rsd1)
```

	year		name	obs	hat	residual	residualLag	qqx	qqy	qqHat
12	1967	Japan	LL II	0.25	0.19	0.26	0.28	0.69	0.26	0.27
13	1968	Japan	LL II	0.26	0.19	0.28	0.38	0.77	0.28	0.29
14	1969	Japan	LL II	0.27	0.19	0.38	0.12	1.57	0.38	0.57
15	1970	Japan	LL II	0.21	0.18	0.12	0.13	0.24	0.12	0.11
16	1971	Japan	LL II	0.21	0.18	0.13	-0.20	0.36	0.13	0.15
17	1972	Japan	LL II	0.14	0.18	-0.20	-0.30	-0.62	-0.20	-0.19

The columns identify the observations (`year`, `name` and may include other covariates such as age, season, etc.), the original observations (`obs`) and the fitted values and the residuals (`obs`, `hat`) if **diags** has been used to read in the data, the residuals with a lag of 1 (`residualLag`) and the quantiles (`qqx`, `qqy`, `qqHat`) assuming a normal distribution.

In some assessment packages the data are in a specific file in other cases the data are in a suite of files found in a dir. Therefore the `readCpue` takes either a file or a dir as its first argument depending on the assessment method e.g. reading in from `vpa2box` and `SS`

Creating a scatter plot is therefore very similar to how it is done using base R, i.e.

```
> u2box=readCpue("unisex09.c01", "2box")
> uSS =readCpue("myDir", "ss")
```

`readCpue` only reads in the data as input to a stock assessment, `diags` reads the residuals and covariates as well.

```
> u2box=readCpue("unisex09.c01", "2box")
> uSS =readCpue("myDir", "ss")
```

There is also the `writeCpue` method for writing the various input files

### 3. Transformations

For plotting and analysis the data may need to be transformed, e.g. observations scaled so that they can be compared, or Pearson residuals computed. This can be done as required using `transform` and `plyr`, e.g. to standardise the residuals or scale them so that they lie between 0 and 1.

```
> stdz( rnorm(10,1,.3))
> minMax(rnorm(10,1,.3))
```

If you wish to scale the residuals within a series then the `plyr` can be used which implements the split-apply-combine strategy for R e.g.

```
> rsdl=ddply(rsdl, .(name), transform, stdRsdl=stdz(residual))
```

One common definition, known as the Pearson residual, is as follows, however the definition depends on the law of large numbers, so it works less well where the number of points in each

Table 1: plot

Syntax	
x	vector values of one variable of interest
data	<code>data.frame</code> other variables
geom	<code>ggplot</code> object that sets the type of plot to construct, i.e. <code>point</code> , <code>line</code> , <code>histogram</code> .

series is relatively small. Therefore in this analysis we used the raw residuals, but the residuals could be transformed as required.

There may be other analyses that are useful to run on the raw data, e.g. running a GAM to calculate a common trend that the individual series can be compared to. I.e. to look for trends that may be different from the others. This can be done by fitting a smoother to year and a series effect, the latter scales the series so that they can be compared, e.g.

```
> library(gam)
> gm =gam(log(obs)~lo(year)+name,data=rsdl)
> rsdl=data.frame(rsdl,gam=predict(gm),gamRsdl=residuals(gm))
> scl =coefficients(gm)[3:9]
> names(scl)=substr(names(scl),5,nchar(names(scl)))
> rsdl=transform(rsdl,scl=scl[as.character(name)])
> rsdl[is.na(rsdl$scl),"scl"]=0
```

## 4. Plotting

Plotting is done using [ggplot2](#), a plotting package based on the grammar of graphics. For those already familiar with lattice a comparison can be found at [ggplot v lattice](#) and [learning R](#). Searching for ggplot2 in google images will throw up lots of examples with code and there is a very active .

### 4.1. Grammar of Graphics

When producing a graphic you have to map data to the visual properties of geometric shapes (e.g. points, lines areas). This may require statistical transformations of the data, a coordinate system that positions the geometric objects on the page and facetting where multiple plots can be generated. Each of these tasks are independent and the grammar breaks these into four. First we load up FLR and an example data set based on North Sea plaice. ggplot uses data in the form of a data.frame so we next have to convert the FLR object to a data.frame.

Facetting creates individual panels by the facetting variables, while themes allow you to prettify the plots.

## 5. Exploratory Data Analysis

First the CPUE time series are plotted using `geom_line` to plot the common trend as estimated by the GAM, then `geom_smooth` fits a loess by series and then `geom_point` is used to overlay the original observations. `facet_wrap` then plots the series individually. `theme_ms` is a bespoke theme to change the look of the plot from the default.

The correlations between indices can be seen by plotting the indices against each other

The indices are then be grouped based on a cluster analysis

To explore the groupings further series within a cluster are plotted together

```

> ggplot(rsd1)+ geom_line(aes(year,exp(gam)),col="red") +
+   geom_smooth(aes(year,obs),se=FALSE) +
+   geom_point(aes(year,obs,col=name)) +
+   facet_wrap(~name,ncol=1,scale="free_y") +
+   theme_ms(legend.position="none") +
+   xlab("Year") + ylab("Index")

```




Figure 1: Plot of indices of abundance, points are the observed index values and the blue a lowess fit to the points by index. The red line is GAM fitted to  $\log(\text{year})$  and fleet.

```

> uMat=ddply(rsd1,.(name),transform, obs=stdz(obs))
> uMat=cast(uMat,year~name,value="obs")
> uMat=uMat[apply(uMat,1,function(x) !all(is.na(x))),]
> pM=plotmatrix(uMat[, -c(1:2,4)])
> pM$layers[[2]]=NULL
> mns=ddply(subset(pM$data,! (is.na(x) & !is.na(y))),.(xvar,yvar), function(x) mean(x$y,na.
> pM+geom_hline(aes(yintercept=V1),data=mns,col="red") +
+   geom_smooth(method="lm",fill="blue", alpha=0.1) +
+   theme(legend.position="bottom") +
+   xlab("Index")+ylab("Index") +
+   theme_ms()

```




Figure 2: Pairwise scatter plots of the indices of abundance, blue lines are linear regressions fitted to the points, the shade area is the standard error of predicted means and the red line is the mean of the points on the y-axis.

```

> cr=cor(uMat[,-1],use="pairwise.complete.obs")
> dimnames(cr)=list(gsub("_"," ",names(uMat)[-1]),gsub("_"," ",names(uMat)[-1]))
> cr[is.na(cr)]=0
> corrplot(cr,diag=F,order="hclust",addrect=2) +
+       theme(legend.position="bottom") +
+       theme_ms()

```

NULL




Figure 3: A plot of the correlation matrix for the indices, blue indicate a positive correlation and red negative. the order of the indices and the rectangular boxes are chosen based on a hierarchical cluster analysis using a set of dissimilarities for the indices being clustered.



```

> ggplot(subset(rsd1,name %in% c("Japan LL II","Japan LL III","South Africa BB II")))+
+   geom_point(aes(year,exp((gam+gamRsd1)-scl),col=name))+
+   geom_smooth(aes(year,exp(gam+gamRsd1-scl),group=name,col=name),se=T,fill="b
+   theme_ms(legend.position="bottom") +
+   xlab("Year") +ylab("Index")

```




Figure 4: Plots of Japan LL II, Japan LL III and South Africa BB II indices, lowess regressions and SEs are shown for each series.

Figure 5: Plots of South Africa BB I, South Africa BB II and Uruguay LL indices, lowess regressions and SEs are shown for each series.

```

> ggplot(subset(rsd1,name %in% c("South Africa BB I","South Africa BB II","Uruguay LL")
+       geom_point(aes(year,exp((gam+gamRsd1)-scl),col=name))+
+       geom_smooth(aes(year,exp(gam+gamRsd1-scl),group=name,col=name),se=T,fill="b
+       theme_ms(legend.position="bottom") +
+       xlab("Year") +ylab("Index")

```




Figure 6: Observed CPUE verses fitted, blue line is a linear regression fitted to points, black the  $y=x$  line.

## 6. Residual Analysis

Next the fit to the indices is evaluated by plotting the residuals. The first plot is of the observed and the predicted values. Since  $U = qB$ , i.e. the index is assumed to be proportional to stock size the points should fall either side of the  $y = x$  line.

Departures from the assumption that the index is proportional to the stock can also be seen by plotting the residuals by time.

Autocorrelated residuals may mean that the estimated parameters are biased, autocorrelation can be checked by plotting the residuals against each other with a lag e.g.

The error distribution can be checked by plotting the observed and the predicted quantiles for a given distribution e.g. for the normal distribution

The variance

```

> dat=ddply(rsd1, .(name), with, data.frame(obs=stdz(obs),hat=stdz(hat)))
> ggplot(dat) +
+   geom_abline(aes(0,1))
+   geom_point(aes(obs,hat))
+   stat_smooth(aes(obs,hat),method="lm",fill="blue", alpha=0.1)
+   facet_wrap(~name,ncol=3)
+   theme_ms(legend.position="bottom")
+   xlab("Fitted") + ylab("Observed")

```

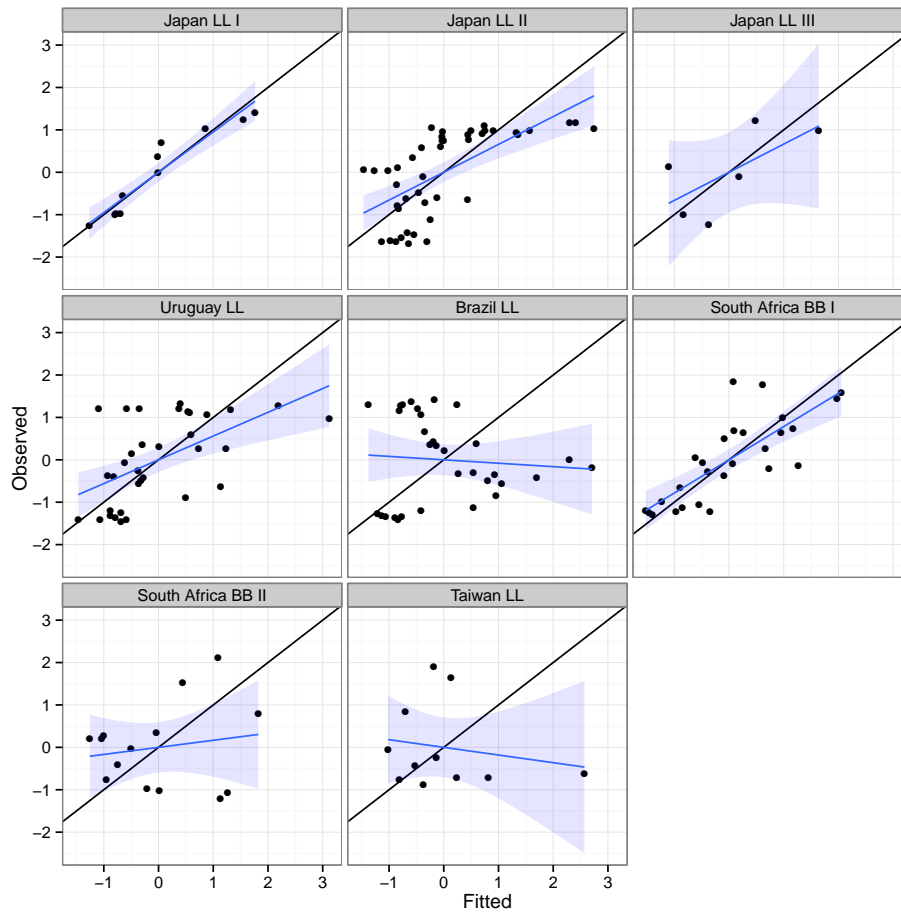


Figure 7: Observed CPUE verses fitted, blue line is a linear regression fitted to points, black the  $y=x$  line.

```

> dat=ddply(rsd1, .(name), transform, residual=stdz(residual,na.rm=T))
> ggplot(aes(year,residual),data=dat) +
+   geom_hline(aes(yintercept=0))      +
+   geom_point()                      +
+   stat_smooth(method="loess",se=T,fill="blue", alpha=0.1) +
+   facet_wrap(~name,scale="free",ncol=2) +
+   theme_ms()

```




Figure 8: Residuals by year, with lowess smoother and SEs.

```

> ggplot(rsd1)
+   geom_point(aes(residual,residualLag))
+   stat_smooth(aes(residual,residualLag),method="lm",se=T,fill="blue", alpha=0.1)
+   geom_hline(aes(yintercept=0))
+   facet_wrap(~name,scale="free",ncol=3)
+   xlab(expression(Residual[t])) +
+   ylab(expression(Residual[t+1])) +
+   theme_ms(legend.position="bottom")

```

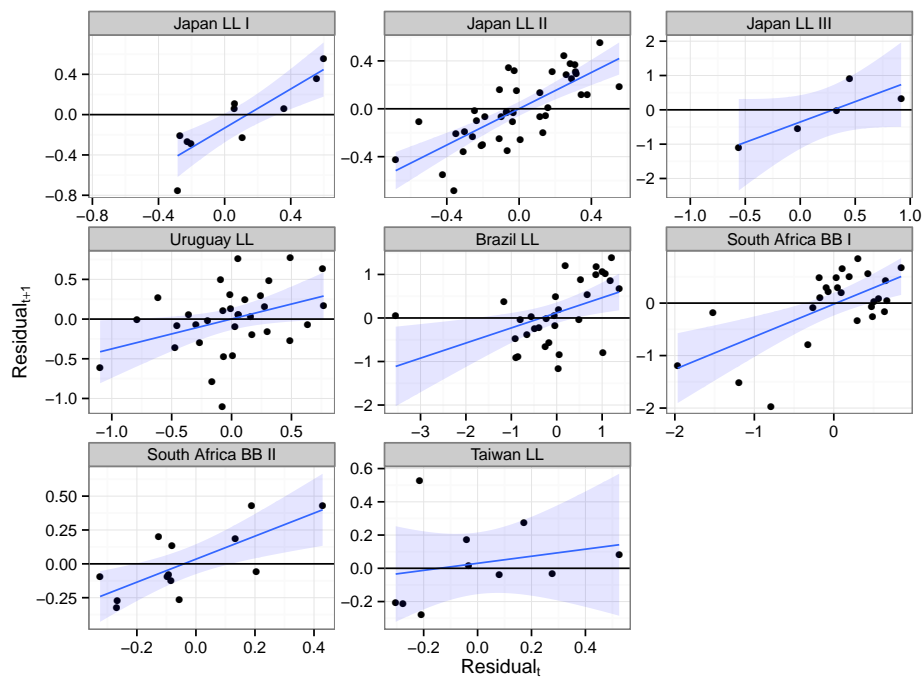


Figure 9: **Plot of autocorrelation, i.e.  $residual_{t+1}$  verses  $residual_t$ .**

```

> ggplot(rsd1)
+   geom_point( aes(qqx,qqy))
+   stat_smooth(aes(qqx,qqHat),method="lm",se=T,fill="blue", alpha=0.1)
+   facet_wrap(~name)
+   theme_ms(legend.position="bottom")
+   theme_ms()

```




Figure 10: Quantile-quantile plot to compare residual distribution with the normal distribution.

```

> ggplot(aes(hat, residual), data=rsdl) +
+   geom_hline(aes(yintercept=0)) +
+   geom_point() +
+   stat_smooth(method="loess", span=.9, fill="blue", alpha=0.1) +
+   facet_wrap(~name, scale="free", ncol=3) +
+   theme_ms()

```




Figure 11: Plot of residuals against fitted value, to check variance relationship.



## 7. Standardised CPUE

Most CPUE series used in stock assessment have been standardised using a Generalised Linear Model (GLM). This requires choosing an appropriate error distribution, variance function and link function ?.

The best way to check these assumptions are by plotting, best performed for a model that included all the main factors (i.e the most-complex model) since if the most complex model isn't a reasonable fit, then any simpler models that are selected will fit adequately because if they didn't they wouldn't be selected.

Going clockwise from the top left in figure 12 the first panel is a q-q plot to check that the residuals follow a normal distribution, the standardised deviance residuals are then plotted against the fitted values to check for systematic departures from the assumptions underlying the error distribution, then the absolute values of the residuals against the fitted values as a check of the assumed variance function and finally the dependent variable against the linear predictor function as a check of the assumed link function ?.

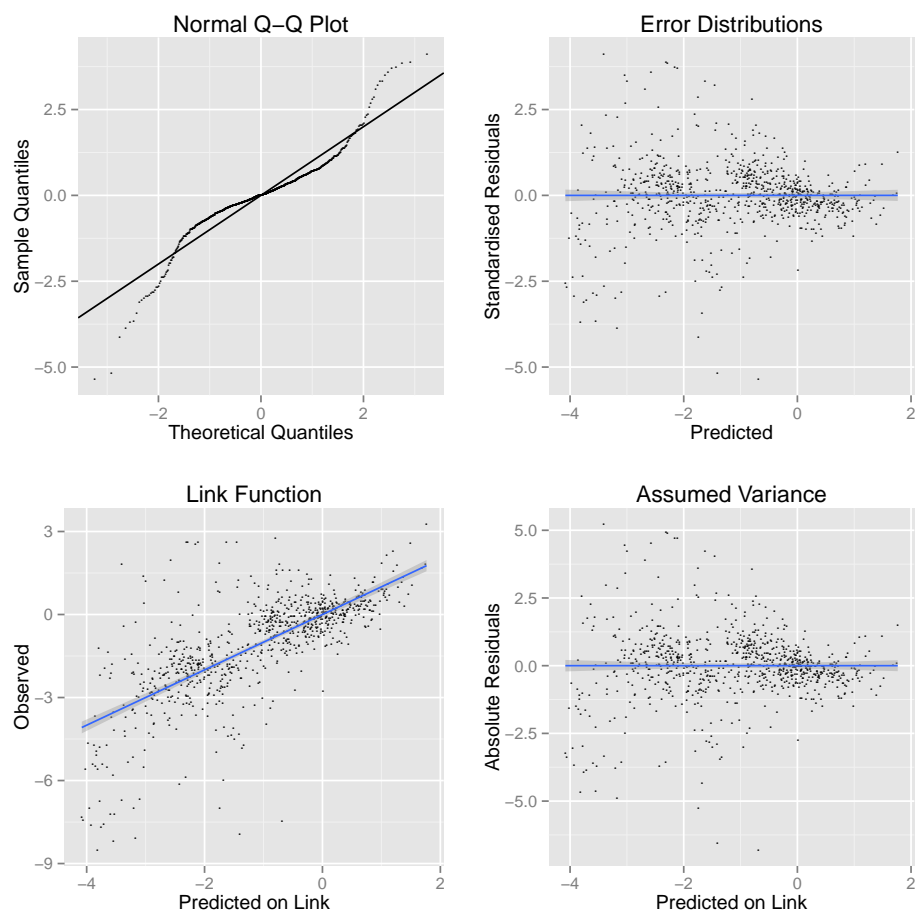


Figure 12: Plot of residuals against fitted value, to check variance relationship.

**Affiliation:**

Laurence Kell  
ICCAT Secretariat  
C/Corazón de María, 8.  
28002 Madrid  
Spain

E-mail: [Laurie.Kell@iccat.int](mailto:Laurie.Kell@iccat.int)